

## **CLOUD COMPUTING LAB 3**

### **1. Why is the Auto Scaling feature significant in cloud computing?**

Auto scaling is a cloud computing feature that automatically adjusts the amount of computational resources allocated to an application based on its current demand. This means that the system can dynamically increase or decrease the number of active instances of services, virtual machines, or containers to maintain optimal performance and efficiency.

The Auto Scaling feature is significant in cloud computing for several key reasons:

#### **1. Cost Efficiency:**

- Optimized Resource Utilization: Auto scaling ensures that you are using only the necessary resources at any given time, scaling up during peak demands and scaling down during low demands. This prevents over-provisioning and reduces costs associated with idle resources.
- Pay-as-you-go Model: By dynamically adjusting the resources based on current needs, users can take full advantage of the pay-as-you-go pricing model offered by cloud providers, leading to significant cost savings.

#### **2. Improved Performance and Availability:**

- Handling Variable Loads: Auto scaling can respond to unexpected spikes in demand, ensuring that applications remain responsive and perform well under varying loads.
- High Availability: By automatically adding or removing instances based on demand, auto scaling helps maintain high availability and reliability of applications, minimizing the risk of downtime.

#### **3. Enhanced Flexibility and Scalability:**

- Elasticity: Auto scaling provides the elasticity needed to adapt to changing workloads without manual intervention. This is crucial for applications with unpredictable or highly variable usage patterns.
- Scalability: Businesses can seamlessly grow their infrastructure as needed, supporting business growth and new user requirements without significant upfront investments in hardware.

#### **4. Operational Efficiency:**

- Reduced Manual Intervention: Automating the scaling process reduces the need for manual monitoring and intervention, freeing up IT teams to focus on other critical tasks.

- Automated Management: Auto scaling can manage the entire lifecycle of instances, from provisioning to termination, ensuring that the infrastructure is always optimally configured.

## **5. Enhanced User Experience:**

- Consistent Performance: By maintaining optimal resource levels, auto scaling helps ensure a consistent and positive user experience, even during peak usage times.
- Responsiveness: Quick response to changing demands means that users experience minimal latency and disruptions.

## **6. Disaster Recovery and Fault Tolerance:**

- Resilience: Auto scaling can help create more resilient architectures by automatically replacing failed instances and distributing the load among healthy instances.
- Disaster Recovery: In the event of a failure, auto scaling can quickly provision new instances, helping to maintain service continuity.

## **2. Describe the following**

- 1. Metric-based autoscaling**
- 2. Schedule-based autoscaling**

### **1. Metric Based Auto Scaling**

Metric Based autoscaling, also known as dynamic or reactive autoscaling, automatically adjusts the number of active instances of a service or resource based on real time performance metrics. This ensures that the infrastructure adapts to varying workloads, providing sufficient resources during high demand and scaling down during low demand. Key metrics often used for this type of autoscaling include:

- CPU Utilization: Scaling up or down based on the percentage of CPU usage.
- Memory Usage: Adjusting instances according to memory consumption.
- Network Traffic: Scaling in response to the volume of inbound or outbound network traffic.
- Custom Metrics: Application Specific metrics such as request latency, error rates, or any other user defined parameters.

Metric Based auto scaling provides a responsive and efficient way to manage resources, ensuring optimal performance and cost efficiency by only using the necessary resources at any given time.

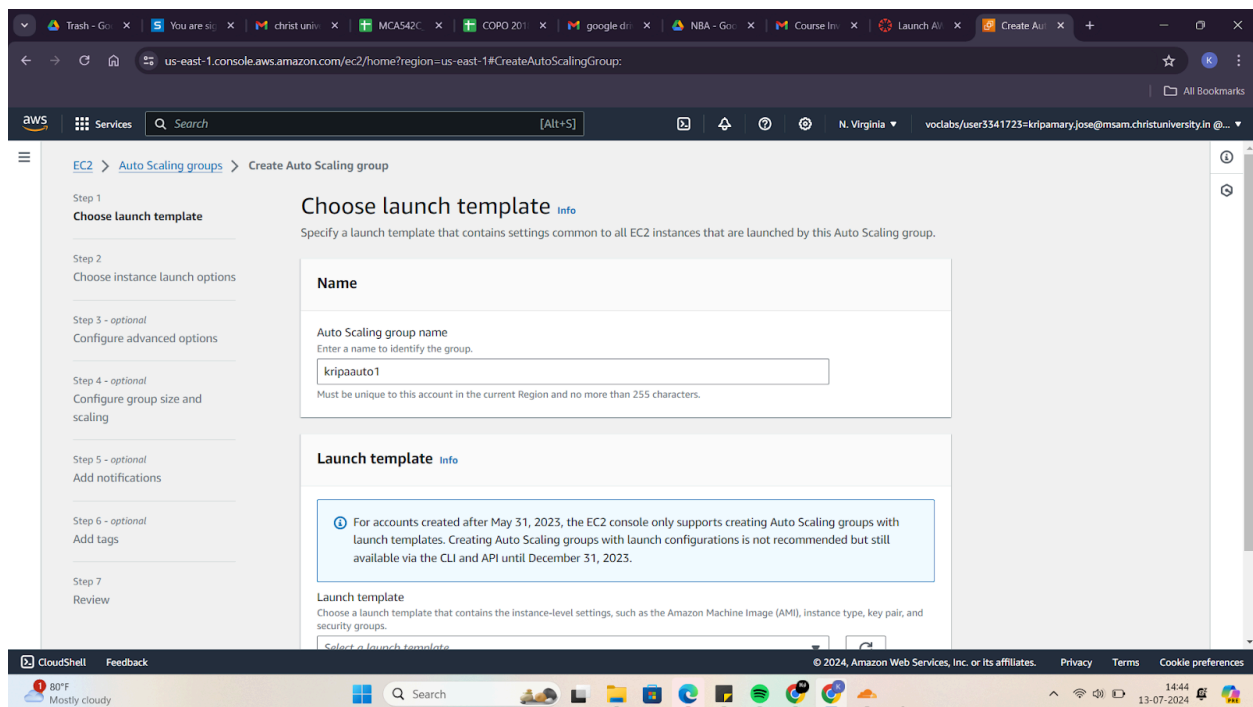
## 2. Schedule Based Auto Scaling

Schedule Based auto scaling, also known as time based or predictive autoscaling, adjusts the number of instances based on predefined schedules rather than real time metrics. This type of auto scaling is ideal for predictable workloads where resource needs are known to fluctuate at specific times. Examples of scenarios where schedules based auto scaling is useful include:

- Business Hours: Increasing resources during peak business hours and scaling down after hours or on weekends.
- Seasonal Trends: Adjusting resources for known seasonal spikes, such as holiday shopping periods.
- Regular Maintenance: Scaling down during scheduled maintenance windows or scaling up in anticipation of planned events.

Schedule Based auto scaling ensures that the necessary resources are available during known periods of high demand and helps to optimize costs by reducing resources during known periods of low activity.

## 3. Demonstrate Metric based auto scaling or Schedule based auto scaling to cater your organizations business requirements.





us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateAutoScalingGroup:

aws

Services

Search

[Alt+S]

N. Virginia

voclabs/user3341723=kripamary.jose@msam.christuniversity.in @...

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1  
Choose launch template

Step 2  
Choose instance launch options

Step 3 - optional  
Configure advanced options

Step 4 - optional  
Configure group size and scaling

Step 5 - optional  
Add notifications

Step 6 - optional  
Add tags

Step 7  
Review

Choose instance launch options

info

You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.

Override launch template

Launch template	Version	Description
kriptemp1	Default	-
lt-0af4fea45d9708b34		
Instance type		
t2.micro		

Network

info

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

80°F

Mostly cloudy

Search

14:46

13-07-2024

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateAutoScalingGroup:

aws

Services

Search

[Alt+S]

N. Virginia

voclabs/user3341723=kripamary.jose@msam.christuniversity.in @...

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1  
Choose launch template

Step 2  
Choose instance launch options

Step 3 - optional  
Configure advanced options

Step 4 - optional  
Configure group size and scaling

Step 5 - optional  
Add notifications

Step 6 - optional  
Add tags

Step 7  
Review

Configure advanced options - optional

info

Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.

Load balancing

info

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer

☐ Attach to an existing load balancer

☒ Attach to a new load balancer

Attach to a new load balancer

info

Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type

Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, visit the Load Balancing console.

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

80°F

Mostly cloudy

Search

14:47

13-07-2024

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#AutoScalingGroupDetails&id=kripaauto1;view=details

Services

Search

[Alt+S]

N. Virginia

voclabs/user3341723=kripamary.jose@msam.christuniversity.in@...

EC2 > Auto Scaling groups > kripaauto1

kripaauto1

DetailsActivityAutomatic scalingInstance managementMonitoringInstance refresh

Group details

Auto Scaling group name  
kripaauto1

Desired capacity  
2

Minimum capacity  
1

Maximum capacity  
3

Date created  
Sat Jul 13 2024 14:48:01 GMT+0530 (India Standard Time)

Desired capacity type  
Units (number of instances)

Status  
-

Amazon Resource Name (ARN)  
arn:aws:autoscaling:us-east-1:777047694729:autoScalingGroup:14a0f3ce-5385-4c34-bb1f-5e0a170555a:autoScalingGroupName/kripaauto1

Launch template

Launch template  
lt-0af4fea45d9708b34

AMI ID  
ami-0b72821e2f351e396

Instance type  
t2.micro

Owner  
arn:aws:sts::777047694729:assumed-

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

Weather Tomorrow

Search

14:48  
13-07-2024

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#targetGroup:targetGroupArn=arn:aws:elasticloadbalancing:us-east-1:777047694729:targetgroup/kripatarg1/1912ee15582c87c0

Services

Search

[Alt+S]

N. Virginia

voclabs/user3341723=kripamary.jose@msam.christuniversity.in@...

EC2 DashboardEC2 Global ViewEventsConsole-to-CodePreview

Instances

InstancesInstance TypesLaunch TemplatesSpot RequestsSavings PlansReserved InstancesDedicated HostsCapacity Reservations

Images

AMIsAMI Catalog

Elastic Block Store

Volumes

2  
Total targets

2  
Healthy

0  
Unhealthy

0  
Unused

0  
Initial

0  
Draining

0 Anomalous

Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

TargetsMonitoringHealth checksAttributesTags

Registered targets (2) Info

Anomaly mitigation: Not applicableDeregisterRegister targets

Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

Filter targets

	Instance ID	Name	Port	Zone	Health status	Health status details	Launch...
<input type="checkbox"/>	i-0c286b5b8d1b96c34		80	us-east-1f	Healthy	-	July 13, 2...
<input type="checkbox"/>	i-0219e9cdef4c13142		80	us-east-1e	Healthy	-	July 13, 2...

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

80°F  
Mostly cloudy

Search

14:51  
13-07-2024

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#LoadBalancer:loadBalancerArn=arn:aws:elasticloadbalancing:us-east-1:777047694729:loadbalancer/app/kripaload1/Saeb2ac6e7ad...

aws Services Search [Alt+S] N. Virginia voclabs/user3341723=kripamary.jose@msam.christuniversity.in @...

EC2 Dashboard  
EC2 Global View  
Events  
Console-to-Code [Preview](#)

▼ Instances  
Instances  
Instance Types  
Launch Templates  
Spot Requests  
Savings Plans  
Reserved Instances  
Dedicated Hosts  
Capacity Reservations

▼ Images  
AMIs  
AMI Catalog

▼ Elastic Block Store  
Volumes

EC2 > Load balancers > kripaload1

kripaload1 Refresh Actions

▼ Details

Load balancer type Application	Status Active	VPC <a href="#">vpc-09a86a921fdb729c4</a>	Load balancer IP address type IPv4
Scheme Internet-facing	Hosted zone Z35SXDOTRQ7X7K	Availability Zones <a href="#">subnet-09432380de8444b59</a> us-east-1f (use1-az5) <a href="#">subnet-0df09ee2f0d7ffc22</a> us-east-1d (use1-az4) <a href="#">subnet-0e019ec7638237694</a> us-east-1c (use1-az2) <a href="#">subnet-05192d2af57288240</a> us-east-1a (use1-az6) <a href="#">subnet-0a048568f117b10b1</a> us-east-1b (use1-az1) <a href="#">subnet-005457db2bd4084ed</a> us-east-1e (use1-az3)	Date created July 13, 2024, 14:48 (UTC+05:30)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 80°F Mostly cloudy 14:51 13-07-2024

Not secure kripaload1-1752452913.us-east-1.elb.amazonaws.com

Hostname: ip-172-31-71-167.ec2.internal

80°F Mostly cloudy 14:52 13-07-2024